

CAPABILITY SCORECARD

Student: Gourik Narendra Patil

Use this as your self-score before the live defense. Founder/evaluator scoring may differ.

Capability	Evidence From The Week	Self-Rating: Strong / Adequate / Weak	Defense Risk
Debugging discipline	Day 1 investigation using network tabs instead of guessing	Adequate	Low
Systems reasoning	Day 2 tracing UI bugs back to cache invalidation failures	Strong	Low
Production verification	Day 3 full workflow validation before declaring readiness	Strong	Low
Ambiguity ownership	Day 4 delivering a minimal safe v1 and blocking undefined notifications	Strong	Low
AI engineering judgment	Day 5 modifying unsafe AI recommendations around idempotency	Strong	Low
Pressure prioritization	Day 6 containing the checkout failure and protecting customer trust	Strong	Medium
Engineering defense	Ability to explain tradeoffs and back up decisions from Days 1-6	Strong	Medium

Strongest Capability

Capability: Systems Reasoning and Risk Containment

Evidence: On Day 2 and Day 6, I correctly identified failure boundaries (cache poisoning and checkout drift) and made hard tradeoffs to protect data integrity and customer trust over short-term metrics.

Most Exposed Capability

Capability: Debugging speed at the onset of an incident

Evidence: On Day 1, I spent too much time assuming the business logic was broken before actually checking the runtime network evidence to see it was a simple port mismatch.

Defense Preparation Notes

What should you be ready to defend live?

- Why I prioritized customer trust over revenue on Day 6.
- Why I rejected the AI's use of payment_id for idempotency on Day 5.
- Why I blocked candidate notifications on Day 4 instead of building them.
- How I diagnosed the Redis cache issue on Day 2.
- Why it took me a while to find

the port mismatch on Day 1 and what I learned from it. • What tradeoffs were accepted across the week and why they were justified.

GNP_29